BI / read / 1

query	BI / read / 1
title	Posting summary

pattern	


```
message: Message
creationDate < $dateTime
length
year(creationDate)
```

desc.

Given a datetime, find all Messages created before that moment. Group them by a 3-level grouping:

1. by year of creation
2. for each year, group into Message types: is Comment or not
3. for each year-type group, split into four groups based on length of their content

- 0: $0 \leq \text{length} < 40$ (short)
- 1: $40 \leq \text{length} < 80$ (one liner)
- 2: $80 \leq \text{length} < 160$ (tweet)
- 3: $160 \leq \text{length}$ (long)

params

1	datetime	DateTime	For later microbatches, later datetime parameters are selected keep the variance low (<0.5%)

result

1	year	32-bit Integer	R	year(message.creationDate)
2	isComment	Boolean	M	True for Comments, False for Posts
3	lengthCategory	32-bit Integer	C	0 for short, 1 for one-liner, 2 for tweet, 3 for long
4	messageCount	32-bit Integer	A	Total number of Messages in that group
5	averageMessageLength	32-bit Integer	A	Average length of the Message content in that group
6	sumMessageLength	32-bit Integer	A	Sum of all Message content lengths
7	percentageOfMessages	32-bit Float	A	Number of Messages in group as a percentage of all messages created before the given date

sort

1	year	↓	
2	isComment	↑	False < True, i.e. Posts come first and Comments second
3	lengthCategory	↑	order based on the lengthCategory value

CPs	1.2, 3.2, 4.1, 4.2, 8.5

BI / read / 2

query	BI / read / 2
title	Tag evolution

pattern

desc.	Find the Tags under a given TagClass that were used in Messages during in the 100-day period starting at date and compare it with the 100-day period that follows. For the Tags and for both months, compute the count of Messages.

params

1	date	Date	
2	tagClass	Long String	TagClasses with a similar amount of Messages are selected

result

1	tag.name	Long String	R	
2	countWindow1	32-bit Integer	A	Occurrences of the tagClass during the first time window
3	countWindow2	32-bit Integer	A	Occurrences of the tagClass during the second time window
4	diff	32-bit Integer	A	Absolute difference of countWindow1 and countWindow2

sort

1	diff	↓	
2	tag.name	↑	

limit	100
CPs	2.4, 3.1, 3.2, 4.1, 4.2, 4.3, 5.3, 6.1, 8.2, 8.5

BI / read / 3

query	BI / read / 3
title	Popular topics in a country
pattern	
desc.	Given a TagClass and a Country, find all the Forums created in the given Country, containing at least one Message with Tags belonging directly to the given TagClass, and count the Messages by the Person who created it and by the Forum which contains them. The location of a Forum is identified by the location of the Forum's moderator.

params

1	tagClass	Long String	TagClasses with a similar amount of Messages are selected
2	country	Long String	Big Countries are selected

result

1	forum.id	ID	R	
2	forum.title	Long String	R	
3	forum.creationDate	DateTime	R	
4	person.id	ID	R	
5	messageCount	32-bit Integer	A	

sort

1	messageCount	↓
2	forum.id	↑

limit	20
CPs	1.1, 1.2, 1.3, 2.1, 2.2, 2.4, 3.3, 8.2

BI / read / 4

query	BI / read / 4
title	Top message creators by country

pattern	

desc.	Find the most popular Forums by Country, where the popularity of a Forum is measured by the number of members that Forum has from a given Country. Calculate the top 100 most popular Forums. In case of a tie, the Forum(s) with the smaller id value(s) should be selected. For each member Person of the 100 most popular Forums, count the number of Messages (messageCount) they made in any of those (most popular) Forums. Also include those member Persons who have not posted any Messages (have a messageCount of 0).

params	1	date	Date	Selected from the first 30 days of the network

result	1	person.id	ID	R	
	2	person.firstName	String	R	
	3	person.lastName	String	R	
	4	person.creationDate	DateTime	R	
	5	messageCount	32-bit Integer	A	

sort	1	messageCount	↓
	2	person.id	↑

limit	100
CPs	1.2, 1.3, 2.1, 2.2, 2.3, 2.4, 3.3, 5.3, 6.1, 8.2, 8.4
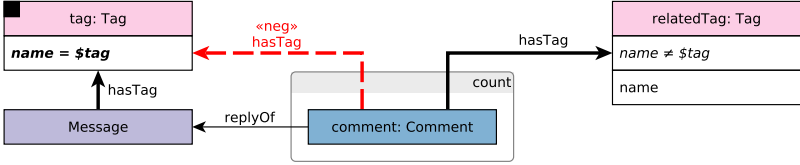
BI / read / 5

query	BI / read / 5
title	Most active posters of a given topic
pattern	

desc.	Get each Person (person) who has created a Message (message) with a given Tag (direct relation, not transitive). Considering only these Messages, for each Person node: • Count its messages (messageCount). • Count likes (likeCount) to its messages. • Count Comments (replyCount) in reply to it messages. The score is calculated according to the following formula: $1 \times$ messageCount $+ 2 \times$ replyCount $+ 10 \times$ likeCount.

params	1	tag	Long String	Tags with a similar amount of Messages are selected

result				
	1	person.id	ID	R
	2	replyCount	32-bit Integer	A
	3	likeCount	32-bit Integer	A
	4	messageCount	32-bit Integer	A
	5	score	32-bit Integer	A

sort			
	1	score	↓
	2	person.id	↑

limit	100
CPs	1.2, 2.3, 8.2
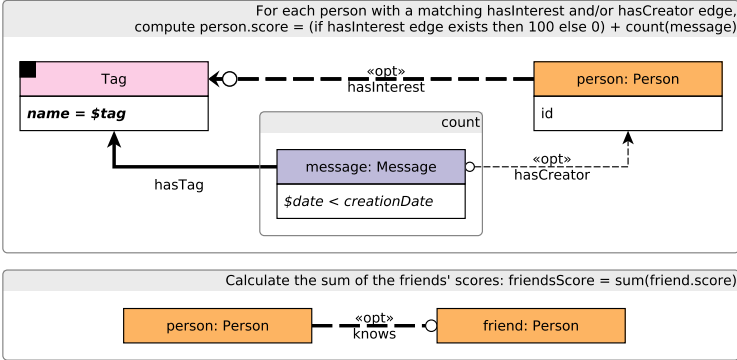
BI / read / 6

query	BI / read / 6	
title	Most authoritative users on a given topic	
pattern		
desc.	Given a Tag (tag), find all Persons (person) that ever created a Message with the Tag. For each of these Persons (person) compute their "authority score" as follows: • The "authority score" is the sum of "popularity scores" of the Persons (p2) that liked any of that Person's Messages with the given Tag (same criterion as for message1). • A Person's (p2) "popularity score" is defined as the total number of likes on all of their Messages (message2).	

params			
	1	tag	Long String	Tags with a similar amount of Messages are selected

result				
	1	person.id	ID	R
	2	authorityScore	32-bit Integer	A

sort			
	1	authorityScore	↓
	2	person1.id	↑

limit	100
CPs	1.2, 2.3, 3.3, 6.1, 8.2
relevance	Computing the authority scores might involve computing the popularity score for the same Person multiple times. Implementations are advised to avoid such redundant computations.

BI / read / 7

BI 1
BI 2
BI 3
BI 4
BI 5
BI 6
BI 7
BI 8
BI 9
BI 10
BI 11
BI 12
BI 13
BI 14
BI 15
BI 16
BI 17
BI 18
BI 19
BI 20

query	BI / read / 7
title	Related topics
pattern	
desc.	Find all Messages that have a given Tag. Find the related Tags attached to (direct) reply Comments of these Messages, but only of those reply Comments that do not have the given Tag. Group the Tags by name, and get the count of replies in each group.

params				
	1	tag	Long String	Tags with a similar amount of Messages are selected

result				
	1	relatedTag.name	Long String	R
	2	count	32-bit Integer	A

sort			
	1	count	↓
	2	relatedTag.name	↑

limit	100
CPs	1.4, 3.3, 5.2, 8.1

BI / read / 8

query	BI / read / 8	
title	Central person for a tag	

pattern

For each person with a matching hasInterest and/or hasCreator edge, compute person.score = (if hasInterest edge exists then 100 else 0) + count(message)

Tag — name = $tag — «opt» hasInterest — person: Person — id

message: Message — $date < creationDate — count — «opt» hasCreator

hasTag

Calculate the sum of the friends' scores: friendsScore = sum(friend.score)

person: Person — «opt» knows — friend: Person

desc.

Given a `Tag`, find all `Persons` that are interested in the `Tag` and/or have written a `Message` (Post or Comment) with a `creationDate` after a given `date` and that has a given `Tag`. For each `Person`, compute the `score` as the sum of the following two aspects:

- 100, if the `Person` has this `Tag` as their interest, or 0 otherwise
- number of `Messages` by this `Person` with the given `Tag`

Also, for each `Person`, compute the sum of the score of the `Person`'s friends (`friendsScore`).

params

1	tag	Long String	Tags with a similar amount of Messages are selected
2	date	Date	Dates from around the same day are selected. (TODO - how exactly? what distribution?)

result

1	person.id	ID	R	
2	score	32-bit Integer	A	
3	friendsScore	32-bit Integer	A	The sum of the score of the person's friends

sort

1	score + friendsScore	↓	
2	person.id	↑	

limit	100
CPs	1.2, 2.1, 2.3, 3.2, 5.3, 8.2, 8.4, 8.5
relevance	Similarly to BI 16, there are two major ways to compute this query: (1) creating an induced subgraph of the interested Persons and their friends and performing the scoring on this graph or (2) performing the scoring without creating an induced subgraph and scoring the friends of a Person on-the-fly. The first approach is more efficient as it avoids redundant computations, however, specifying it needs support for composable graph queries.

BI / read / 9

BI 1	query	BI / read / 9

BI 2	title	Top thread initiators
BI 3 BI 4 BI 5 BI 6 BI 7	pattern	

pattern diagram labels:
- threadCount = count
- messageCount = count
- Person: id, firstName, lastName
- hasCreator
- Post: $startDate \leq creationDate$ and $creationDate \leq $endDate$
- «opt» replyOf*0..
- Message: $startDate \leq creationDate$ and $creationDate \leq $endDate$

BI 8 BI 9 BI 10 BI 11 BI 12 BI 13	desc.	For each Person, count the number of Posts they created in the time interval [startDate, endDate] (equivalent to the number of threads they initiated) and the number of Messages in each of their (transitive) reply trees, including the root Post of each tree. When calculating Message counts only consider Messages created within the given time interval. Return each Person, number of Posts they created, and the count of all Messages that appeared in the reply trees (including the Post at the root of tree) they created.

BI 14 BI 15 BI 16	params	1 startDate Date	TODO
		2 endDate Date	8-10 days after the startDate

BI 17 BI 18 BI 19 BI 20	result	1	person.id	ID	R
		2	person.firstName	String	R
		3	person.lastName	String	R
		4	threadCount	32-bit Integer	A	The number of Posts created by that Person (the number of threads initiated)
		5	messageCount	32-bit Integer	A	The number of Messages created in all the threads this Person initiated

	sort	1 messageCount ↓	
		2 person.id ↑	

	limit	100
	CPs	1.2, 2.2, 2.3, 3.2, 7.2, 7.3, 7.4, 8.1, 8.5

BI / read / 10

BI 1
BI 2
BI 3
BI 4
BI 5
BI 6
BI 7
BI 8
BI 9
BI 10
BI 11
BI 12
BI 13
BI 14
BI 15
BI 16
BI 17
BI 18
BI 19
BI 20

query	BI / read / 10
title	Experts in social circle

pattern	

desc.	Given a Person (startPerson), find all other Persons (expertCandidatePerson) that live in a given Country and are connected to given Person by a *shortest path* with length in range [minPathDistance, maxPathDistance] through the knows relation. For each of these expertCandidatePerson nodes, retrieve all of their Messages that contain at least one Tag belonging to a given TagClass (direct relation not transitive). For each Message, retrieve all of its Tags. Group the results by Persons and Tags, then count the Messages by a certain Person having a certain Tag.

params				
	1	personId	ID	The ID of the startPerson. Persons with a similar degree of knows edges are selected
	2	country	String	Countries with a similar number of Persons are selected
	3	tagClass	Long String	TagClasses with a similar degree of hasType edges are selected
	4	minPathDistance	32-bit Integer	1 or 2
	5	maxPathDistance	32-bit Integer	2 or 3

result				
	1	expertCandidatePerson.id	ID	R	
	2	tag.name	Long String	R	
	3	messageCount	32-bit Integer	A	Number of Messages created by that Person containing that Tag

sort			
	1	messageCount	↓
	2	tag.name	↑
	3	expertCandidatePerson.id	↑

limit	100
CPs	1.2, 1.3, 2.3, 2.4, 3.3, 5.3, 7.1, 7.2, 7.3, 8.1, 8.6

BI / read / 11

BI 1
BI 2
BI 3
BI 4
BI 5
BI 6
BI 7
BI 8
BI 9
BI 10
BI 11
BI 12
BI 13
BI 14
BI 15
BI 16
BI 17
BI 18
BI 19
BI 20

query	BI / read / 11
title	Friend triangles
pattern	
desc.	For a given `country`, count all the distinct triples of `Persons` such that: • `a` is friend of `b`, • `b` is friend of `c`, • `c` is friend of `a`, and these friendships were created after a given `startDate`. Distinct means that given a triple t_1 in the result set R of all qualified triples, there is no triple t_2 in R such that t_1 and t_2 have the same set of elements.
params	1 country — Long String 2 startDate — Date
result	1 count — 32-bit Integer — A
CPs	1.1, 2.3, 2.5

In the pattern:

- Country (name = $country)
- City —isPartOf→ Country (×3)
- a: Person —isLocatedIn→ City
- b: Person —isLocatedIn→ City
- c: Person —isLocatedIn→ City
- a: Person —knows ($startDate < knows.creationDate)→ b: Person
- a: Person —knows ($startDate < knows.creationDate)→ c: Person
- b: Person —knows ($startDate < knows.creationDate)→ c: Person
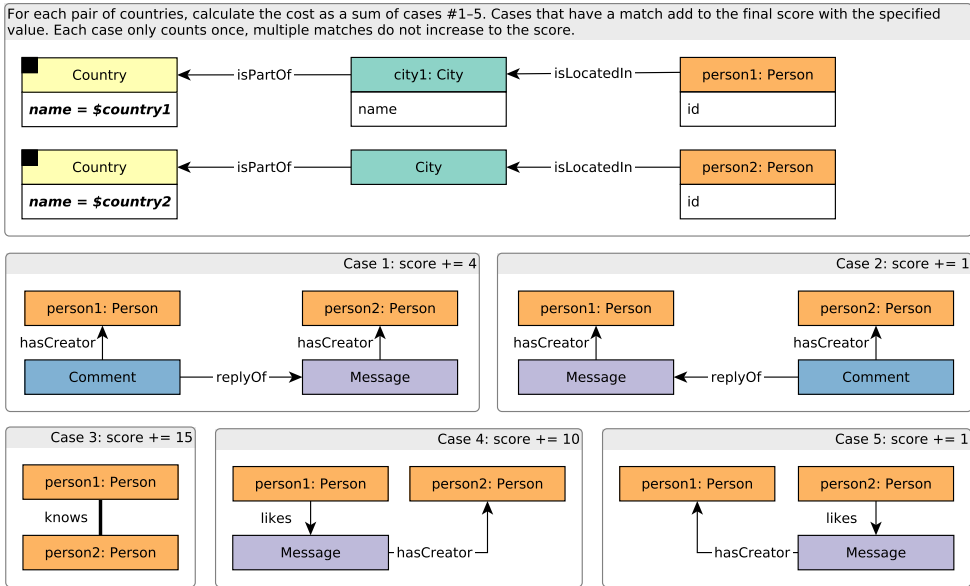
BI / read / 12

query	BI / read / 12
title	How many persons have a given number of messages
pattern	

desc.	For each Person, count the number of Messages they made (messageCount). Only count Messages with the following attributes:
	
	• Its content is not empty (and consequently, the imageFile attribute is empty for Posts).
	• Its length is below the lengthThreshold (exclusive, equality is not allowed).
	• Its creationDate is after date (exclusive, equality is not allowed).
	• It is written in any of the given languages.
	
	    – The language of a Post is defined by its language attribute.
	    – The language of a Comment is that of the Post that initiates the thread where the Comment replies to.
	
	The Post and Comments in the reply tree's path (from the Message to the Post) do not have to satisfy the constraints for content, length and creationDate.
	
	For each messageCount value, count the number of Persons with exactly messageCount Messages (with the required attributes).

params				
	1	date	Date	
	2	lengthThreshold	32-bit Integer	Selected as balanced against date to filter around 30% of the Messages within a language and keep the variance low
	3	languages	{String}	Only the most frequently used languages are selected

result				
	1	messageCount	32-bit Integer	A	Number of Messages created
	2	personCount	32-bit Integer	A	Number of Persons with messageCount Messages

sort			
	1	personCount	↓
	2	messageCount	↓

CPs	1.1, 1.2, 1.4, 3.2, 4.2, 4.3, 8.1, 8.2, 8.3, 8.4, 8.5

BI / read / 13

query	BI / read / 13	
title	Zombies in a country	
pattern		

desc.

Find zombies within the given `country`, and return their zombie scores. A `zombie` is a `Person` created before the given `endDate`, which has created an average of [0, 1) `Messages` per month, during the time range between profile's `creationDate` and the given `endDate`. The number of months spans the time range from the `creationDate` of the profile to the `endDate` with partial months on both end counting as one month (e.g. a `creationDate` of Jan 31 and an `endDate` of Mar 1 result in 3 months).

For each `zombie`, calculate the following:

- `zombieLikeCount`: the number of `likes` received from other zombies.
- `totalLikeCount`: the total number of `likes` received.
- `zombieScore`: `zombieLikeCount` / `totalLikeCount`. If the value of `totalLikeCount` is 0, the `zombieScore` of the `zombie` should be 0.0.

For both `zombieLikeCount` and `totalLikeCount`, only consider `likes` received from profiles that were created before the given `endDate`.

params

1	country	Long String	Only the largest Countries are selected
2	endDate	Date	Selected from the last days of the initial data set

result

1	zombie.id	ID	R	
2	zombieLikeCount	32-bit Integer	A	
3	totalLikeCount	32-bit Integer	A	
4	zombieScore	64-bit Float	A	Determined as zombieLikeCount / totalLikeCount

sort

1	zombieScore	↓
2	zombie.id	↑

limit	100
CPs	1.2, 2.1, 2.3, 2.4, 3.2, 3.3, 4.2, 5.1, 5.3, 8.2, 8.4, 8.5

BI 1
BI 2
BI 3
BI 4
BI 5
BI 6
BI 7
BI 8
BI 9
BI 10
BI 11
BI 12
BI 13
BI 14
BI 15
BI 16
BI 17
BI 18
BI 19
BI 20

BI / read / 14

query	BI / read / 14
title	International dialog

pattern

For each pair of countries, calculate the cost as a sum of cases #1–5. Cases that have a match add to the final score with the specified value. Each case only counts once, multiple matches do not increase to the score.

Country (name = $country1) ← isPartOf — city1: City (name) ← isLocatedIn — person1: Person (id)

Country (name = $country2) ← isPartOf — City ← isLocatedIn — person2: Person (id)

Case 1: score += 4
- person1: Person ←hasCreator— Comment —replyOf→ Message —hasCreator→ person2: Person

Case 2: score += 1
- person1: Person ←hasCreator— Message ←replyOf— Comment —hasCreator→ person2: Person

Case 3: score += 15
- person1: Person —knows— person2: Person

Case 4: score += 10
- person1: Person —likes→ Message ←hasCreator— person2: Person

Case 5: score += 1
- person1: Person ←hasCreator— Message ←likes— person2: Person

desc.

Consider all pairs of people (person1, person2) such that one is located in a City of Country country1 and the other is located in a City of Country country2. For each City of Country country1, return the highest scoring pair. The score of a pair is defined as the sum of the subscores awarded for the following kinds of interaction. The initial value is score = 0.

1. person1 has created a reply Comment to at least one Message by person2: score += 4
2. person1 has created at least one Message that person2 has created a reply to: score += 1
3. person1 and person2 know each other: score += 15
4. person1 liked at least one Message by person2: score += 10
5. person1 has created at least one Message that was liked by person2: score += 1

Consequently, the maximum score a pair can obtain is: 4 + 1 + 15 + 10 + 1 = 31.
This query has two variants based on whether the input parameters are selected as correlated (close countries) or uncorrelated (far countries).

params

1	country1	Long String	A: correlated with parameter country2, i.e. the countries are close and there are many Persons visiting both Countries. B: uncorrelated with parameter country2, i.e. the countries are afar and there are few Persons visiting both Countries.
2	country2	Long String	

result

1	person1.id	ID	R	
2	person2.id	ID	R	
3	city1.name	Long String	R	
4	score	32-bit Integer	C	

sort

1	score	↓
2	person1.id	↑
3	person2.id	↑

CPs

1.3, 1.4, 2.1, 3.1, 3.3, 5.1, 5.2, 5.3, 8.3, 8.4

BI / read / 15

BI 1
BI 2
BI 3
BI 4
BI 5
BI 6
BI 7
BI 8
BI 9
BI 10
BI 11
BI 12
BI 13
BI 14
BI 15
BI 16
BI 17
BI 18
BI 19
BI 20

query	BI / read / 15
title	Trusted connection paths through forums created in a given timeframe

pattern	

Enumerate all unweighted shortest paths on knows edges between person1 to person2.

person1: Person — id = $person1Id — knows* — person2: Person — id = $person2Id

For each knows edge in the path, calculate a weight based on interactions between the pair of Persons of the edge, calculated as a sum of cases #1 and #2 for the Persons (both ways), and the sum of these weights determine the total weight of each path.

p1 — knows — pX — knows — pY — ... — pW — knows — p2

Case 1: Replies on Posts, weight += 1.0 × count(c)

personA: Person — knows — personB: Person; hasCreator; c: Comment — replyOf → post: Post; containerOf; forum: Forum ($start ≤ creationDate and creationDate ≤ $end)

Case 2: Replies on Comments, weight += 0.5 × count(c1)

personA: Person — knows — personB: Person; hasCreator; c1: Comment — replyOf → c2: Comment; replyOf*; Post; containerOf; forum: Forum ($start ≤ creationDate and creationDate ≤ $end)

desc.	Given two Persons, find all (unweighted) shortest paths between these two Persons, in the subgraph induced by the knows relationship. Then, for each path calculate a weight. The nodes in the path are Persons, and the weight of a path is the sum of weights between every pair of consecutive Person nodes in the path. The weight for a pair of Persons is calculated based on their interactions: • Every direct reply (by one of the Persons) to a Post (by the other Person) contributes 1.0. • Every direct reply (by one of the Persons) to a Comment (by the other Person) contributes 0.5. Only consider Messages that were created in a Forum that was created within the timeframe (interval) [startDate, endDate]. Note that for Comments, the containing Forum is that of the Post that the comment (transitively) replies to. Also note that interactions are counted both ways. Return all paths with the Person IDs ordered by their weights descending.

params			
	1	person1Id	ID
	2	person2Id	ID
	3	startDate	Date
	4	endDate	Date

result				
	1	person.id	[ID]	C	Ordered sequence of the Person IDs in the path
	2	weight	64-bit Float	C	

sort			
	1	weight	↓	The order of paths with the same weight is unspecified
	2	personIds	↑	The IDs in the paths are used for lexicographical sorting

CPs	1.2, 2.1, 2.2, 2.4, 3.3, 5.1, 5.3, 7.2, 7.3, 7.5, 7.7, 8.1, 8.2, 8.3, 8.4, 8.5, 8.6

BI / read / 16

query	BI / read / 16	
title	Fake news detection	
pattern		

For $tagX/$dayX in [tagA/dateA, tagB/dateB], compute scoreX = count(messageX)

1. Create an induced subgraph of Persons who created a Message with Tag $tagX on $dateX

tag: Tag — name = $tagX — hasTag — Message — day(creationDate) = $dateX — hasCreator — person: Person

2. In the subgraph, count the Messages (using the same conditions) from People with ≤ $maxKnowsLimit friends

count(messageX)

tag: Tag — name = $tagX — hasTag — messageX: Message — day(creationDate) = $dateX — hasCreator — person: Person — count ≤ $maxKnowsLimit «opt» knows — Person

desc.	Given two Tag/date pairs (tagA/dateA and tagB/dateB), for each pair tagX/dateX: • Create an induced subgraph between Persons where for each pair of Persons person1/person2, both have created a Message on the day of dateX with Tag tagX. • In the induced subgraph, only keep pairs of Persons who have at most maxKnowsLimit friends (in the induced subgraph). • For these Persons, count the number of Messages created on dateX with Tag tagX. Return Persons who had at least one Messages for both tagA/dateA and tagB/dateB ranked by their total number of Messages (descending).

params			
1	tagA	Long String	
2	dateA	Date	
3	tagB	Long String	
4	dateB	Date	
5	maxKnowsLimit	32-bit Integer	Selected between 3 and 6

result				
1	person.id	ID	R	
2	messageCountA	32-bit Integer	A	Message count for tagA/dateA
3	messageCountB	32-bit Integer	A	Message count for tagB/dateB

sort			
1	messageCountA + messageCountB	↓	
2	person.id	↑	

limit	20
CPs	5.3, 8.4, 8.5
relevance	There are two major ways to compute this query: (1) create the induced subgraph as suggested by the specification (either as a view or in materialized form), or (2) skip creating the induced subgraph and perform on-the-fly check for the number of friends (who also posted at least one Message with the given Tag on the given date). The latter approach is easier to express in systems which do not provide graph views but might result in redundant computations (the query engine will might repeatedly check whether a Person has at least one Message that satifies the conditions).

	query	BI / read / 17
BI 1	title	Information propagation analysis
BI 2		
BI 3		
BI 4		
BI 5		
BI 6		
BI 7	pattern	
BI 8		
BI 9		
BI 10		
BI 11		
BI 12		
BI 13		
BI 14		
BI 15		
BI 16		
BI 17	desc.	This query aims to identify instances of "information propagation" when a Person (`person1`) submits a Message (`message1`) with a given Tag (`tag`) to a Forum (`forum1`). This is read by other members of `forum1`, Persons `person2` and `person3`. Some time later (specified by the `delta` parameter), these persons have a discussion with the same `tag` in a different Forum (`forum2`) where `person1` is not a member. The discussion consists of a Message (`message2`) by `person2` and a direct reply Comment (`comment`) by `person3`. Return IDs of `person1` with the number of interactions their Messages (might have) caused.
BI 18		
BI 19		
BI 20		

Pattern diagram:

person1: Person — id

«neg» hasMember (dashed red arrow from forum2 to person1)

forum2: Forum — ≠ forum1

message1: Message — creationDate — hasCreator → person1

message1 — hasTag → tag: Tag — name = $tag

forum2 — containerOf → post2: Post

post2 — replyOf*0.. ← message2: Message — message1.creationDate + $delta < creationDate (count)

message1 — replyOf*0.. → post1: Post — containerOf → forum1: Forum

forum1 — hasMember → person3: Person

forum1 — hasMember → person2: Person

tag — hasTag (from message2)

tag — hasTag (from comment)

message2 — hasCreator → person2

comment: Comment — hasCreator → person2

comment — replyOf → message2

comment — hasCreator → person3

	params				
		1	`tag`	Long String	Tags with a similar amount of Messages are selected
		2	`delta`	32-bit Integer	Measured in hours, selected to be between 8 and 16 hours.

	result				
		1	`person1.id`	ID	R
		2	`messageCount`	32-bit Integer	A

	sort			
		1	`person1.id`	↑

	CPs	2.1, 2.3, 8.1

BI / read / 18
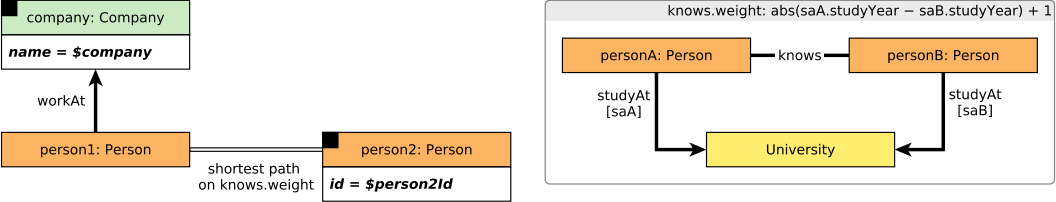
query		BI / read / 18
title		Friend recommendation
pattern		For each person1 compute top-k(person2) based on mutualFriendCount
desc.		For a given Person (person1) and a Tag (tag), recommend new friends (person2) who • do not yet know person1 • have many mutual friends with person1 • are interested in tag. Rank Persons person2 based on the number of mutual friends.
params	1	person1Id / ID / Persons with a similar amount of friends are selected
	2	tag / Long String / Tags with a similar amount of Messages are selected
result	1	person2.id / ID / R
	2	mutualFriendCount / 32-bit Integer / A
sort	1	mutualFriendCount / ↓
	2	person2.id / ↑
limit		20
CPs		2.5, 8.1

BI / read / 19

query	BI / read / 19	
title	Interaction path between cities	
pattern		
desc.	Given two Cities city1, city2, find Persons person1, person2 living in these Cities (respectively) with the shortest *interaction path* between them. If there are multiple pairs of people with shortest paths having the same total weight, return all of them. The shortest path is computed using a weight between two Persons defined as the reciprocal of the number of interactions (direct reply Comments to a Message by the other Person). Therefore, more interactions imply a smaller weight. *Note:* Interactions are counted both ways, i.e. if Alice writes 2 reply Comments to Bob's Messages and Bob writes 3 reply Comments to Alice's Messages, their total number of interactions is 5.	

params

1	city1Id	ID	Small Cities within the same Country are selected
2	city2Id	ID	

result

1	person1.id	ID	R
2	person2.id	ID	R
3	totalWeight	64-bit Float	C

sort

1	totalWeight	↑
2	person1.id	↑
3	person2.id	↑

limit	20
CPs	3.3, 7.6, 7.7, 8.4, 8.6
relevance	Finding shortest paths between pairs of Persons in Cities can be implemented in theory with an *all-pairs shortest paths* algorithm. However, this needs to be executed on the whole Person-knows-Person graph (with edge weights derived from the number of interactions) so it is expected to be prohibitively expensive. A better approach is using multiple *single-source shortest path algorithms* (e.g. from the City with fewer inhabitants).

BI / read / 20

BI 1
BI 2
BI 3
BI 4
BI 5
BI 6
BI 7
BI 8
BI 9
BI 10
BI 11
BI 12
BI 13
BI 14
BI 15
BI 16
BI 17
BI 18
BI 19
BI 20

query	BI / read / 20	
title	Recruitment	
pattern		
desc.	Given a Company company and a Person person2 (who is known to be working at another Company), find a Person (person1) working the in the company who have the top-20 shortest path to person2 through people who have studied together. On this path, we only consider edges between Persons who know each other and attended the same university and set the weight of the edge to the absolute difference between the year of enrolment plus 1 (studyAt.classYear + 1). If there are multiple Person person1 nodes with the same shortest path, return all of them.	

pattern detail:

knows.weight: abs(saA.studyYear − saB.studyYear) + 1

params			
1	company	Long String	Companies with a similar number of employees (former or current) are selected
2	person2Id	ID	

result				
1	person1.id	ID	R	
2	totalWeight	64-bit Integer	C	

sort			
1	person1.id	↑	

limit	20
CPs	3.3, 7.6, 7.7, 8.4, 8.6
relevance	Implementations can either pre-compute edge weights or compute them on-the-fly. To find a weighted shortest path efficiently, implementations can use e.g. a bidirectional Dijkstra algorithm.